

```
In [58]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
from sklearn.linear_model import LinearRegression
from sklearn.neighbors import KNeighborsRegressor
warnings.filterwarnings('ignore')
```

Part 1: Data Exploration and Visualization (20 marks)

Q1 loading

```
In [9]: df = pd.read_csv('lifesat.csv');
df.head()
```

```
Out[9]:   Country  GDP per capita (USD)  Life satisfaction
0    Russia      26456.387938          5.8
1    Greece      27287.083401          5.4
2    Turkey      28384.987785          5.5
3    Latvia      29932.493910          5.9
4   Hungary      31007.768407          5.6
```

Q2


```
In [14]: print(df.info())
print(df.describe())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 27 entries, 0 to 26
Data columns (total 3 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Country          27 non-null    object  
 1   GDP per capita (USD) 27 non-null    float64 
 2   Life satisfaction  27 non-null    float64 
dtypes: float64(2), object(1)
memory usage: 780.0+ bytes
None
GDP per capita (USD)  Life satisfaction
count                27.000000      27.000000
mean                 41564.521771     6.566667
std                  9631.452319     0.765607
min                 26456.387938     5.400000
25%                 33938.289305     5.900000
50%                 41627.129269     6.800000
75%                 49690.580269     7.300000
max                 60235.728492     7.600000
```

Q3

```
In [41]: plt.figure(figsize=(10, 6))
sns.scatterplot(df, x="GDP per capita (USD)", y="Life satisfaction", alpha=0.7).set_title('GDP for Life satisfaction')
plt.show()
```

GDP for Life satisfaction in Countries



There is a positive correlation

Part 2: Linear Regression Model (30 marks)

Q4

```
In [29]: X = df[['GDP per capita (USD)']]
y = df[['Life satisfaction']]
print(X.shape, y.shape, sep='\n')

(27, 1)
(27, 1)
```

Q5


```
In [46]: lr = LinearRegression().fit(X, y)
print("Coefficient: ", lr.coef_[0])
print("Intercept: ", lr.intercept_)

Coefficient:  [6.77889969e-05]
Intercept:  [3.74904943]
```

Q6

```
In [52]: plt.figure(figsize=(10, 6))
# plt.scatter(X, y, color='blue', label='Data')
sns.scatterplot(df, x="GDP per capita (USD)", y="Life satisfaction", alpha=0.7).set_title('GDP for Life satisfaction')
plt.plot(X, lr.predict(X), color='red', linewidth=3, label='Linear Prediction')
plt.legend()
plt.show()
```

GDP for Life satisfaction in Countries



Q7

```
In [75]: prediction_lr = lr.predict([[37655.2]])
print(prediction_lr[0])
```

[6.30165767]

Part 3: K-Nearest Neighbors Regression (25 marks)

```
In [59]: knn = KNeighborsRegressor(n_neighbors=3)
knn.fit(X, y)
```

```
Out[59]:
```

▼ KNeighborsRegressor ⓘ ⓘ

▼ Parameters

⌚	n_neighbors	3
⌚	weights	'uniform'
⌚	algorithm	'auto'
⌚	leaf_size	30
⌚	p	2
⌚	metric	'minkowski'
⌚	metric_params	None
⌚	n_jobs	None

Q9

```
In [67]: prediction_knn = knn.predict([[37655.2]])
prediction_lr = lr.predict([[37655.2]])
print(prediction_knn[0], prediction_lr[0], sep='\n')
```

```
[6.33333333]
[6.30165767]
```

Q10

```
In [77]: ks = [1, 3, 5, 10]
ans = []
for k in ks:
    model = KNeighborsRegressor(n_neighbors=k)
    model.fit(X, y)
    pred = model.predict([[37655.2]])[0]
    ans.append(pred)

print(ans)
plt.figure(figsize=(8, 5))
plt.plot(ks, ans, marker='o', linestyle='--')
plt.title("KNN Predictions for different k values")
plt.xlabel("Number of Neighbors")
plt.ylabel("Predicted Life satisfaction")
plt.xticks(ks)
plt.show()
```

```
[array([7.2]), array([6.33333333]), array([6.26]), array([6.37])]
```

